

Cloud Native Observability.

A horizontally scalable, highly available backend for ingesting telemetry (metrics + logs), storing time-series data, firing alerts, and serving queries to a Realtime dashboard.

High-level architecture (backend-focused)

- **Ingress / API Gateway**
 - TLS termination, rate limit, auth, routing to services
- **Ingestion Layer**
 - **Ingest API** (HTTP/gRPC) accepts metrics & logs
 - **Buffering:** Kafka (or Redis Streams) between ingestors and writers
 - **Ingestors:** stateless workers that validate, batch, and forward to storage
- **Persistence**
 - **TimescaleDB** (Postgres + Timescale) for metrics (hypertables + continuous aggregates)
 - **Postgres (vanilla)** for metadata, users, alert rules, dashboards
 - **Elasticsearch or ClickHouse** for high-volume logs & full-text search (optional/hybrid)
- **Processing**
 - **Aggregator Service:** builds rollups / aggregates, writes to continuous aggregates or materialized views
 - **Alerting Service:** subscribes to aggregated metrics, evaluates rules, triggers notifications
- **Query API / Backend for Dashboard**
 - Query endpoints for timeseries, logs, metadata; supports pagination, filters, bucketed aggregation
 - WebSocket / SSE for real-time push (alerts, live metrics)
- **Auth & User Service**
 - JWT + refresh tokens / API keys for services and users
 - Role-based access control (RBAC)
- **Admin / Management**
 - Schema migrations, telemetry of the platform itself
- **Observability of the observability**
 - Prometheus metrics + traces (OTel) for all services

Technology choices (recommended)

- Language: **Go** or **.NET (C#)** or **Node (TypeScript)** for microservices (Go preferred for lightweight concurrency & shipping single binaries)
- Message broker: **Kafka** (production) / **Redis Streams** (simpler)
- Metrics DB: **TimescaleDB** (Postgres + Timescale)
- Logs DB (optional): **Elasticsearch** or **ClickHouse**
- API Gateway: **Traefik** / **NGINX Ingress** on Kubernetes
- Container orchestration: **Kubernetes**
- IaC: **Terraform** (cloud) + **Helm** or kustomize (cluster)
- CI/CD: **GitHub Actions** (or Azure DevOps/GitLab CI)
- Schema migrations: **Flyway** / **golang-migrate** / **EF Core Migrations**
- Secrets: **Vault** / cloud KMS
- Telemetry: **OpenTelemetry** + **Prometheus** + **Grafana**
- Authentication: **OAuth2** / **JWT** (or integrate with identity provider)

Services & responsibilities (detailed)

1. **api-gateway**

- Accepts all public traffic: /v1/ingest, /v1/query, /v1/auth, /v1/logs
- Rate-limits per API key, basic auth checks, TLS

2. **ingest-service**

- HTTP/gRPC endpoints for clients/agents
- Validates payloads, enriches with metadata, writes to Kafka/topic
- Emits internal metrics

3. **ingest-writer(s)**

- Consumer(s) of Kafka topic
- Batch writes to TimescaleDB (metrics) and Elasticsearch/ClickHouse (logs)
- Idempotency and retry handling

4. **aggregator / rollup-service**

- Creates continuous aggregates or writes downsampled series
- Runs scheduled jobs or stream-processing to compute 1m/5m/1h rollups

5. **query-service**

- Exposes REST GraphQL/gRPC endpoints for dashboard queries
- Caches frequent queries (Redis) and supports pagination

6. **alerting-service**

- Continuously evaluates alert rules (pull from DB or subscribe to processed streams)
- Supports rate-limited notifications (email, webhook, Slack)
- Records alert history in metadata DB

7. **auth-service**

- Issues & validates JWTs, issues API keys, rate-limit keys
- RBAC lookups in metadata DB

8. **user-mgmt & config-service**

- Dashboard configs, users, projects, environments, alert rules stored in Postgres

9. **ingest-agent (optional)**

- Lightweight agent to run on VMs/containers to collect system metrics & ship to ingest endpoint

10. **admin-service / operator tools**

- Tasks: migrations, backup restores, cluster management

API design (core endpoints)

Use REST or gRPC. Example REST endpoints:

Auth

- POST /v1/auth/login → returns JWT
- POST /v1/auth/apikey → create API key (service use)

Ingestion

- POST /v1/ingest/metrics
payload: { time, service, instance_id, metric_name, value, labels }
- POST /v1/ingest/logs
payload: { time, service, instance_id, level, message, labels, json }

Query

- GET /v1/query/metrics?service=&metric=&from=&to=&step= → returns time-buckets
- POST /v1/query/logs → search query with filters, pagination
- GET /v1/query/series → list series metadata (for dashboard selectors)

Alerts / Rules

- GET /v1/alerts
- POST /v1/alerts → create rule { name, expression, window, severity, notify }
- GET /v1/alerts/history?service=&from=&to=

Admin

- GET /v1/health → platform health
- GET /v1/metrics/platform → internal metrics

Realtime

- WebSocket endpoint /v1/ws for push updates & alerts

Security:

- All endpoints require API key or bearer token
- Admin endpoints require elevated role

Data models & schemas (key tables)

Timescale metrics table (hypertable)

```
CREATE TABLE metrics (
    time      TIMESTAMPTZ NOT NULL,
    org_id    UUID        NOT NULL,
    service   TEXT        NOT NULL,
    instance_id TEXT,
    metric_name TEXT      NOT NULL,
    value     DOUBLE PRECISION NOT NULL,
    labels    JSONB,
    PRIMARY KEY (time, org_id, service, metric_name, instance_id)
);
```

```
SELECT create_hypertable('metrics', 'time', chunk_time_interval => INTERVAL '1 day');
```

Metadata (Postgres)

```
-- users
```

```
CREATE TABLE users (
    id UUID PRIMARY KEY,
    email TEXT UNIQUE NOT NULL,
    password_hash TEXT,
    role TEXT,
    created_at TIMESTAMPTZ DEFAULT now()
);
```



```
-- api_keys
```

```
CREATE TABLE api_keys ( id UUID PRIMARY KEY, user_id UUID REFERENCES users(id), key TEXT, scopes TEXT[], created_at TIMESTAMPTZ );
```



```
-- alert_rules
```

```
CREATE TABLE alert_rules (
    id UUID PRIMARY KEY,
    org_id UUID,
    name TEXT,
    expression TEXT, -- e.g., "avg(cpu_usage) > 80"
```

```
window_interval INTERVAL,  
severity TEXT,  
notify JSONB,  
created_by UUID,  
enabled BOOLEAN DEFAULT TRUE,  
created_at TIMESTAMPTZ DEFAULT now()  
);
```

-- alert_history

```
CREATE TABLE alert_history ( id UUID PRIMARY KEY, alert_id UUID REFERENCES alert_rules(id),  
triggered_at TIMESTAMPTZ, state TEXT, details JSONB );
```

Logs (Elasticsearch mapping) — or ClickHouse table if using CH

Store logs as JSON with indexed fields: time, service, instance_id, level, message, plus an analyzed message field for full-text.

Ingestion patterns & batching

- Clients should send batched payloads (100–500 datapoints) to reduce overhead.
- Ingest API writes to Kafka topic metrics.incoming and logs.incoming.
- Writers consume batches, convert to COPY/INSERT batches for TimescaleDB (COPY is fastest).
- Use connection pooler (pgbouncer) between writers and Postgres.

Scaling & HA patterns

- **Stateless services:** scale via replica count in K8s HPA (CPU or custom metrics).
- **Kafka:** partition by org_id or service to parallelize writers.
- **Timescale:** use vertical scaling + partitioning; for huge scale, consider multi-node Timescale.
- **Read replicas:** run Postgres replicas for query-service to protect writes.
- **Health checks:** readiness & liveness probes for each pod.
- **Leader election:** for scheduled aggregator jobs (use Kubernetes leader election or distributed lock via Redis).

Durability & retention

- Use chunk retention policies in Timescale (e.g., drop raw 90d, keep rollups longer).
- Compress older chunks (Timescale compression).
- Archive older raw data to S3 (Parquet) for cold storage and restore if needed.
- Snapshot Elasticsearch indices or export ClickHouse.

CI/CD & release strategy (backend)

Repository layout: mono-repo or services per repo. Either OK; recommend mono-repo for a portfolio.

CI pipeline (GitHub Actions)

- push to main:
 1. Lint & unit tests for service changed
 2. Build Docker image (service-specific)
 3. Run integration tests in lightweight k3s / Docker Compose
 4. Publish image to container registry (GitHub Packages / Docker Hub / ACR)
 5. Create Helm chart update / k8s manifests
 6. Deploy to staging via helm (automated)
 7. Run smoke tests
 8. Deploy to production with manual approval (or automated canary)

CD features

- Canary / blue-green deployments (k8s + helm + service mesh optional)
- Rollbacks on health-check failures
- DB migration job in pipeline (reviews + migrations run in pre-prod automatically)

Schema migrations

- Migrations run as Kubernetes Job with migrate binary at deploy time; fail deployment if migration fails.

Observability & testing of the platform

- **Tracing:** OpenTelemetry traces captured for request flows (ingest → writer → DB)
- **Metrics:** Each service exposes Prometheus metrics; configure a Prom server for platform.
- **Logs:** Centralized logs (Elasticsearch/ClickHouse)
- **SLOs & Alerts:** Monitor ingest_latency, write_failures, db_connections, kafka_lag, api_error_rate.

Testing:

- Unit tests per service
- Integration tests: start local TimescaleDB + Kafka + services in GitHub Actions using Docker Compose / testcontainers
- Load tests: k6 or Vegeta to simulate ingestion patterns with burstiness and sustained load
- Chaos tests: kill pods / introduce latency in staging

Security & secrets

- TLS everywhere (ingest endpoints & internal)
- RBAC in Kubernetes
- Use vault / secrets manager for DB passwords, API keys
- Rate-limit unauthenticated or misbehaving clients
- Input validation & schema strictness to avoid injection

Dev & local environment

- docker-compose.yml for local dev:
 - TimescaleDB, Postgres (metadata), Kafka (or Redis), Zookeeper (if Kafka), Elasticsearch (optional), pgbouncer, and local versions of services
- Scripts to seed sample data + example ingestion client

Backup & restore

- Postgres/Timescale: base backups + WAL archiving, test restores
- Elasticsearch / ClickHouse snapshots to object storage
- Regular recovery drills documented in repo

Deliverables & phased roadmap (milestones)

I. Phase 0 — Scaffolding & infra (MVP backend)

- Repo & service skeletons: ingest-service, ingest-writer, query-service, auth-service, alerting-service
- Docker Compose for local dev (TimescaleDB + Kafka / Redis)
- Basic API endpoints: /v1/ingest/metrics, /v1/query/metrics
- Timescale metrics hypertable + basic inserts
- README + architecture diagram
- Deliverable: demo ingest → store → query flow

II. Phase 1 — Core functionality

- Batched writes + COPY optimization
- Kafka buffering + idempotent consumers
- Query API with aggregation + cache (Redis)
- Simple React dashboard integration (real time via WebSockets)
- CI: unit tests + Docker build + push
- Deliverable: functional dashboard showing live metrics for sample services

III. Phase 2 — Alerting & Aggregation

- Continuous aggregates / rollups (1m/5m/1h)
- Alert rule CRUD + evaluator + notification hooks (webhook/email)
- RBAC + API key auth
- Deliverable: alerts triggered by synthetic load, alert history visible in UI

IV. Phase 3 — Scale & Resilience

- Kubernetes manifests / Helm charts
- Production-ready DB config (pgbouncer, backups, replica)
- Autoscaling rules + Kafka partitioning tuning
- Load testing & performance benchmarks
- Deliverable: deploy to cloud cluster (AKS/EKS/GKE) with docs

V. Phase 4 — Logs, Search & Long-term storage

- Integrate Elasticsearch/ClickHouse for logs
- Implement log ingestion & search endpoints
- Cold storage -> S3 + Parquet
- Deliverable: searchable logs & retention policies enforceable

VI. Phase 5 — Harden & polish

- Observability for platform (Prom + Grafana dashboards)
- Security review + secrets management
- Canary deployments, automatic rollbacks
- Production runbook & SLA documentation
- Deliverable: portfolio-ready repo + case study README and diagrams

Repo structure suggestion (mono-repo)

```
/repo-root  
  /docs  
    architecture.md  
    runbooks.md  
  /infra  
    k8s/ (helm charts)  
    terraform/  
    docker-compose.yml  
  /services  
    /ingest-service  
      main.go  
      Dockerfile  
      /migrations  
    /ingest-writer  
    /query-service  
    /alerting-service  
    /auth-service  
  /clients  
    /agent (example agent)  
  /scripts  
    seed-data.sh  
  /ci  
    workflows/  
  README.md
```

Example CI snippet (GitHub Actions - build & test)

```
name: CI

on: [push]

jobs:
  build:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3
      - name: Set up Go
        uses: actions/setup-go@v4
        with: { go-version: '1.20' }
      - name: Build services
        run: |
          cd services/ingest-service && go build -o ingest
          cd ../ingest-writer && go build -o writer
      - name: Run unit tests
        run: |
          go test ./...
      - name: Build Docker images
        run: |
          docker build -t ghcr.io/${{ github.repository }}/ingest-service:pr-${{ github.run_id }} \
          services/ingest-service
      - name: Push images
        uses: docker/login-action@v2
        with:
          registry: ghcr.io
          username: ${{ github.actor }}
          password: ${{ secrets.GHCR_TOKEN }}
```

Metrics & SLO ideas (for platform)

- **Ingest latency (p95) < 200ms**
- **Write success rate > 99.9%**
- **Query latency (p95) < 300ms for 1m window**
- **Kafka consumer lag < 1000 messages**

Security / Compliance notes

- Obfuscate/omit sensitive data from logs
- GDPR: retention & deletion APIs per org id
- Audit logs of admin actions